An appeal is not the same as a remedy

What our computer-checked examples show
Progress report | November 1, 2026

The finding

A person may receive an explanation, have a named responsible officer, and be allowed to appeal an AI-assisted decision - yet the review office may still lack the power to correct it. Our first Lean example checks that these arrangements can coexist under a limited set of written rules.

A simple employment example

Imagine that an AI-assisted hiring decision rejects an applicant and causes harm. A named officer is responsible. The applicant can send an appeal to a desk that can reconsider the decision. In the first version, that desk has no power to correct it. In the second version, we give the desk correction power. Everything else in the example stays the same.

This is a hypothetical case. We have not investigated a real employer or measured harm to actual applicants.

Safeguard in our example	First version	Changed version
Named responsible officer	Yes	Yes
Explanation and human review	Yes	Yes
Accessible appeal	Yes	Yes
Power to correct	No	Yes
Meets responsibility-and-appeal checklist	Yes	Yes
Meets checklist also requiring correction capacity	No	Yes

Why this matters

When assessing an AI appeal process, ask two separate questions: Can the person get a review? Can someone correct the decision if that review finds an error? A positive answer to the first question does not supply the second.

Keeping the promises separate

What the computer checked

Lean checked our proofs about the two versions of the example. The first meets our responsibility-and-appeal checklist but fails the checklist that also requires correction capacity. The second meets both. We also supplied a case that meets the stronger checklist while still having an affected person and a harmful, high-impact decision. The requirements can therefore be met together in this small model.

A further check keeps a duty separate from its fulfillment. A rule can require correction capacity even when the organization does not provide it. The missing safeguard is then a violation of that rule; writing the rule did not make the safeguard appear.

The assumptions behind the result

- The example has one applicant, one decision, one officer, and one appeal desk.

- We stipulate that the decision is consequential, excludes the applicant, and causes harm.

- An accessible desk with review power counts as an appeal channel.

- An accessible desk with correction power counts as correction capacity. This does not promise that every loss can be undone.

- The two checklists are our explicit working interpretations. They represent only a small part of the source document.

How this relates to the encyclical

The source discussion of appeal and responsibility motivates the questions we test. Our choice to represent remedy through correction capacity is an interpretation that must remain open to scrutiny. Lean checks the consequences of our written rules; it cannot decide that our reading is the only correct reading.

The honest limit of this first result

This is a simple example, not yet a major new finding about AI governance. It shows that the method works and that we can expose a missing safeguard without quietly assuming it exists. We have not established that a real AI system is safe, fair, legally compliant, or ethically justified. We have not formalized the whole encyclical.

What comes next

The following pages add public-benefits responsibility, high-stakes timing, and a comparison of narrow and broad appeal duties. The remaining agreed target suite is unfinished. The final project must include an understandable report of its findings, assumptions, and limits. Completing the code alone will not complete the project.

Source: official English Vatican text, Magnifica Humanitas, selected paragraphs 97-111; prototype references 97 and 103-105. Accessed November 1, 2026. Read the source. Technical evidence: project verification logs and docs/TECHNICAL_RESULTS.md.

Who owns this decision?

A named agency and a public contact do not necessarily identify the person assigned to your case.

The example

Imagine an AI-assisted public-benefit denial. The agency accepts responsibility, publishes a contact's name, and assigns every important decision internally. But the contact handles routine inquiries; another employee handles the denial. In the first version, that case officer is not identifiable. Known staffing and internal assignment therefore do not supply an identifiable case owner.

These are stipulated facts in a hypothetical office. We have not audited an actual public-benefits agency.

Question about the denial	Unnamed	Named only	Name + power
Agency and public contact identified?	Yes	Yes	Yes
An officer assigned to the denial?	Yes	Yes	Yes
That officer identifiable?	No	Yes	Yes
That officer has correction power?	Yes	No	Yes
Identifiable case owner with power?	No	No	Yes

What Lean checked

The first version has known staff and internal allocation, yet no identifiable assignee for the denial. The second names that assignee but supplies no correction power. The third provides both. The agency, membership, and assignments stay fixed.

We also checked a sufficient rule: assign the decision to a person and ensure that this person is identifiable. Correction capacity through that person needs a separate authority condition for the same decision.

The practical lesson

Ask who is assigned to this decision, whether that person is identifiable, and whether they can correct it. These are separate questions. A general contact list does not answer all three.

What this still does not establish

Authority does not prove that the officer will act, that the applicant can reach them, or that every loss can be undone. We have not established legal liability, a uniquely correct reading of the source, or major research novelty.

A stop that arrives too late

The same final checklist can hide a crucial difference: whether control reached the system before the harm occurred.

The example

We model a hazardous autonomous action. A stop request is sent, delivered, and followed by a corrected decision record. One run delivers the stop before execution; the other delivers it afterward. Both finish with the stop recorded as delivered and the administrative record corrected. Only the timely run avoids physical loss.

This is a hypothetical case with explicit rules. We assume a delivered stop is obeyed and that none of the five available actions can reverse the physical loss. Those assumptions are not discoveries about real systems.

Check	Timely run	Late run
Stop requested?	Yes	Yes
Stop delivered?	Before execution	After execution
Record corrected afterward?	Yes	Yes
Record-correction officer identifiable?	Yes	Yes
Physical loss in the model?	No	Yes

What Lean checked

The two runs contain the same four events; only the delivery and execution order changes. Lean checked the different outcomes. It also checked that once physical loss occurs, no finite sequence of the available actions removes it. Correcting the record remains possible, but it does not restore the physical outcome.

There is a positive result too: if the stop has already arrived while the state is loss-free, every continuation allowed by these rules stays loss-free. That is a conditional guarantee inside our model, not a proof that real operators always detect danger or respond in time.

The practical lesson

Ask when control became effective, not just whether a human was involved. Also ask what a promised correction can change: the record, the decision, or the physical harm. Those protections are different.

Where the limits matter

The result depends on the stated irreversibility and stop rules. It does not establish that every real loss is permanent or that the encyclical uniquely requires this particular mechanism. We have initial examples in all three agreed case domains; the full research project remains unfinished.

Who gets to appeal?

An internal review does not itself give the affected person an appeal. The duty's starting point matters.

Two explicit readings

The narrow reading requires appeal after exclusion. The broader reading keeps that rule and adds appeal for an affected person in a consequential decision, even before exclusion. These are two working interpretations; we do not claim they exhaust the source's possible readings.

Hypothetical situation	Narrow checklist	Broad checklist
Person excluded; no accessible appeal	Fails	Fails
Consequential decision affects person; no exclusion yet; no appeal	Passes	Fails
Person excluded; accessible appeal provided	Passes	Passes

Named responsibility and internal review stay fixed. 'Passes' means only this limited checklist; neither policy includes a remedy-capacity duty.

What survives both readings

Both reject exclusion without an accessible appeal. This is a checked consequence of the shared rule, not a new moral discovery. Staff review may occur internally while the person has no channel to challenge. After exclusion, that arrangement fails both checklists.

Where the readings differ

An AI-assisted ranking can affect opportunities before exclusion. Without appeal, the narrow checklist can pass and the broader one fails. Its extra guarantee comes from the added duty. Opening person access satisfies both appeal checklists while review and assignment stay fixed.

The practical question

When an organization says 'a human reviews our AI decisions,' ask whether the affected person can challenge, and when that opportunity begins. Review, person access, and the duty's trigger are separate issues.

Assumptions and limits

These are hypothetical employment-style cases. High impact and affectedness are stipulated; actual harm is not assumed. The broader duty is explicit. Appeal alone does not establish remedy capacity, restoration, or full source compliance. The remaining agreed research questions are still unfinished.